

ISYE 6679 Computational Methods

Final Project Report

On

*Parallel Computing in finance for estimating
risk-neutral densities in option prices*

Ana M. Monteiro, Antonio A.F. Santos

Team Members

Karuna Kumar

Rishi Dasgupta

INDEX

Topic	Pg No.
1. Introduction	3
2. Model Outline	3
2.1 Option Pricing	3
2.2 Non-Parametric Risk-Neutral Density	4
2.3 Cross-Validation	5
3. Algorithm Design	7
4. Results	9
5. Inferences and Conclusion	11
6. References	12
7. Appendix	13
7.1 Main Function	13
7.2 GPU Run Function	14
7.3 Kernel Function	18
7.4 Graph Function	19

1. Introduction

This project aims to replicate results from the paper “Parallel Computing in Finance for estimating risk-neutral densities in option prices” by Ana M. Monteiro and Antonio A.F. Santos from the University of Coimbra, Portugal. This paper discusses the computational advantages offered by parallel computing implementations for solving complex calculations over large datasets, which are commonly involved in assessing financial risk.

Financial options markets assume a central role in all economic systems, which makes estimating the options pricing function imperative for generating sound risk-management strategies. For this purpose, non-parametric estimation methods can be used to estimate the options pricing function (as opposed to parametric estimation which may impose no-negligible constraints). However, this requires large datasets and can be computationally intensive, especially in estimating the optimal kernel bandwidths associated with the same.

The paper presents a case study to identify optimal kernel bandwidths for the non-parametric estimation of risk-neutral densities in options prices. For this purpose, two different datasets are used – the intraday S&P500 index and the CBOE’s VIX (volatility index) put/call/strike values extracted from Yahoo Finance between April 16th to April 20th, 2018, for contracts with a maturity date of May 18th, 2018. Two alternate implementations are provided – a sequential C++ program that utilizes CPU on local hardware, and a parallel computing implementation using CUDA/C++ code that utilizes GPU computing on NVIDIA Tesla V100 hardware provided by Pace-ICE.

The algorithm used includes a tailor-made cross-validation criterion function and minimum grid-search methods to identify optimal kernel bandwidths over both datasets. The elapsed time is recorded for grid sizes - 16x16, 32x32, 64x64, 128x128, and 256x256 for both datasets, and the results are compared with its parallel computing implementation.

In the following sections, we will begin by providing a brief outline of the financial model used, including theoretical explanations and motivations for options pricing, risk-neutral density, and the cross-validation function implemented. Then we will discuss the algorithm design for both sequential and parallel computing implementations in C++, followed by a discussion of the computational results from the experiment. Lastly, we will discuss the inferences and implications of the parallel computing approach and future use cases.

2. Model Outline

2.1 Option pricing

A European-type call/put option is a contract giving the right, but not the obligation for the holder to purchase or sell a given stock S , at a predetermined price X , at the maturity date T . A central object in option pricing is the risk-neutral density (RND) which informs derivative security pricing at an expected value at the time of contract maturity. This would allow for lower (neutral) risk in negotiating option contracts as the risk profile of equities can be predicted and assessed prior to maturity.

The expected values of equities at the time of maturity can be calculated through a risk-neutral density, discounted by a risk-free interest rate. Let us consider the price of a call option contract as the function $C(S_t, X, r, \tau, \sigma)$, which depends on the stock price S_t , risk-free interest rate r , time to maturity $\tau = T - t$, and σ , a parameter characterizing the volatility of the stock price. Considering $q(\cdot)$ as the risk-neutral density at maturity, the prices of call and put options are given as follows:

$$C(S_t, X, r, \tau, \sigma) = e^{-r\tau} \int_X^{\infty} (S_T - X)q(S_T) dS_T, \quad (1)$$

$$P(S_t, X, r, \tau, \sigma) = e^{-r\tau} \int_0^X (X - S_T)q(S_T) dS_T. \quad (2)$$

A fundamental result due to [1] and [2] relates the risk-neutral density with the second-order derivative of the option pricing function with respect to the strike. The first and second-order derivatives of (1) and (2) are as follows:

$$\frac{\partial C(X, S_t, r, \tau, \sigma)}{\partial X} \Big|_{X=S_T} = G(S_T | S_t, r, \tau, \sigma) - 1 \quad (3)$$

$$\frac{\partial P(X, S_t, r, \tau, \sigma)}{\partial X} \Big|_{X=S_T} = G(S_T | S_t, r, \tau, \sigma) \quad (4)$$

$$\begin{aligned} \frac{\partial^2 C(X, S_t, r, \tau, \sigma)}{\partial X^2} \Big|_{X=S_T} &= \frac{\partial^2 P(X, S_t, r, \tau, \sigma)}{\partial X^2} \Big|_{X=S_T} \\ &= e^{-r\tau} q(S_T), \end{aligned} \quad (5)$$

Where G is the cumulative distribution function associated with q . These results allow the risk-neutral density to be estimated from the second derivative of the option pricing function. The no-arbitrage assumption implies an additional set of constraints as follows:

$$C(X, S_t, r, \tau, \sigma) \geq 0; \quad P(X, S_t, r, \tau, \sigma) \geq 0, \quad (6)$$

$$-e^{-r\tau} \leq \frac{\partial C}{\partial X}(X, S_t, r, \tau, \sigma) \leq 0; \quad (7)$$

$$0 \leq \frac{\partial P}{\partial X}(X, S_t, r, \tau, \sigma) \leq e^{-r\tau},$$

$$\frac{\partial^2 C}{\partial X^2}(X, S_t, r, \tau, \sigma) \geq 0; \quad \frac{\partial^2 P}{\partial X^2}(X, S_t, r, \tau, \sigma) \geq 0. \quad (8)$$

These constraints provide preliminary information to be considered in the elimination process jointly with sample price data, which allows for a more robust estimation with readily interpretable estimates.

2.2 Non-Parametric Risk-Neutral Density

In the non-parametric setting, observed options prices are subjected to random noise. The theoretical call price $C(S_t, X, r, \tau, \sigma)$ relates to the observed $C^*(S_t, X, r, \tau, \sigma)$ through the regression model as follows:

$$C^*(S_t, X_i, r, \tau, \sigma) = C(S_t, X_i, r, \tau, \sigma) + \varepsilon_i, \quad i = 1, \dots, n_c, \quad (9)$$

Here, n_c is the sample size of all call prices, and the same treatment is given to put prices. The Nadaraya-Watson (NW) estimator is used to approximate $C(S_t, X_i)$ considering a partition for strike values formed by $\{x_j\}_{j=1}^m$. For a given point x_j , the estimates are as follows:

$$\hat{C}(S_t, X_i) = \frac{\sum_{i=1}^{n_c} K\left(\frac{X_i - x_j}{h_c}\right) C^*(S_t, X_i)}{\sum_{i=1}^{n_c} K\left(\frac{X_i - x_j}{h_c}\right)}, \quad j = 1, \dots, m. \quad (10)$$

where $K(\cdot)$ is a kernel function used to weight observations neighboring x_j , which depends on the bandwidth parameter h_c . The fitting to data depends essentially on the bandwidth value, more than the kernel $K(\cdot)$ functional form.

Regarding option contracts observations, call and put prices are available, and both give relevant information for risk-neutral density estimation. By approximating the call (put) option pricing function as \hat{f}_c (\hat{f}_p) Taylor polynomial expansion of order three, at point x_j , $j = 1, \dots, n$.

$$\begin{aligned}\hat{f}_c(x_j) = & \beta_{0,c} + \beta_{1,c}(X_{i,c} - x_j) + \frac{\beta_{2,c}}{2!}(X_{i,c} - x_j)^2 \\ & + \frac{\beta_{3,c}}{3!}(X_{i,c} - x_j)^3,\end{aligned}\quad (11)$$

$$\begin{aligned}\hat{f}_p(x_j) = & \beta_{0,p} + \beta_{1,p}(X_{i,p} - x_j) + \frac{\beta_{2,p}}{2!}(X_{i,p} - x_j)^2 \\ & + \frac{\beta_{3,p}}{3!}(X_{i,p} - x_j)^3,\end{aligned}\quad (12)$$

where $X_{i,c}$ and $X_{i,p}$ are the call and put strikes. The variables $\beta_{0,c}$ and $\beta_{0,p}$ represent the function level estimates for calls and puts, respectively, $\beta_{1,c}$ and $\beta_{1,p}$ are the first derivatives, $\beta_{2,c}$ and $\beta_{2,p}$ are the second derivatives, and finally, $\beta_{3,c}$ and $\beta_{3,p}$ the third derivatives estimates.

Considering the information on the call and put prices and the set of constraints implied from (3) to (8), the risk-neutral density is estimated by solving the problem:

$$\begin{aligned}\text{minimize}_{\beta} \quad & \sum_{i=1}^{n_c} \left(C_i^* - \sum_{k=0}^3 \frac{\beta_{k,c}}{k!} (X_{i,c} - x_j)^k \right)^2 K \left(\frac{X_{i,c} - x_j}{h_c} \right) w_{c,i} \\ & + \sum_{i=1}^{n_p} \left(P_i^* - \sum_{k=0}^3 \frac{\beta_{k,p}}{k!} (X_{i,p} - x_j)^k \right)^2 K \left(\frac{X_{i,p} - x_j}{h_p} \right) w_{p,i} \quad (13)\end{aligned}$$

$$\text{subject to} \quad -\beta_{1,c} + \beta_{1,p} = 1 \quad (14)$$

$$\beta_{2,c} - \beta_{2,p} = 0 \quad (15)$$

$$-e^{-r\tau} < \beta_{1,c} < 0 \quad (16)$$

$$0 < \beta_{1,p} < e^{-r\tau} \quad (17)$$

$$\beta_{0,c}, \beta_{0,p}, \beta_{2,c}, \beta_{2,p} \geq 0, \quad (18)$$

where $\beta = \{\beta_{0,c}, \dots, \beta_{3,c}, \beta_{0,p}, \dots, \beta_{3,p}\}$, and $w_{c,i}$ and $w_{p,i}$ are weights associated with each observation, usually accounting for data heteroscedasticity. In option prices, the weights can be associated with volume or open interest, which usually assume significant values for options near-the-money and tend to zero for deep-out- and deep-in-the-money options.

Solving the problem and finding the estimates for each x_j is not computationally challenging. If the grid associated with x_j contains a few hundred elements, available computational capabilities can solve the problem in a few seconds. The main computational challenge is finding the optimal kernel bandwidths h_c and h_p . For the practical application of these methods, when they are applied sequentially in risk-neutral density updates, excessive computational times needed to calculate bandwidths can turn such time-varying analysis infeasible.

2.3 Cross-Validation

Cross-validation (CV) is a standard method to define the necessary bandwidth values for nonparametric estimators. It is the most computationally intensive element in a nonparametric estimation setting. The common feature is the inclusion in the cross-validation criterion function of a squared-error component. There are many cases where it is sound to

introduce new factors in the CV function apart from the squared-error component, which intends to accommodate the particularities of the problem at hand. By its nature, second derivative estimators are subjected to estimation errors when compared with estimators for the function itself. In a non-adapted framework, the increased estimation errors can lead to estimates incompatible with an interpretation that links them to density functions. Furthermore, the density must not present great shape diversity for different sub-ranges; for example, a highly multi-modal density is difficult to interpret. On the other hand, extreme homogeneity within sub-ranges is also not desirable, for example, a uniform density. Measurements of area, variation, and entropy redefine the cross-validation function to adequately define bandwidth values that influence the shape of estimated density functions.

Regarding standard forms of cross-validation criterion functions, the leave- k_i -out estimator is used instead of the leave-one-out for a fixed design setting. The problem is defined as:

$$\underset{h}{\text{minimize}} \quad \frac{1}{n} \sum_{i=1}^n \frac{1}{k_i} \sum_{j=1}^{k_i} \left(y_{i,j} - \hat{f}_{(-k_i)}(X_i) \right)^2. \quad (19)$$

The area associated with the second derivative function is used to penalize the criterion function when the area associated with the second derivative differs substantially from one,

$$A(h) = 1 + \left| \sum_{j=1}^m (x_j^* - x_{j-1}^*) f^{(2)} \left(\frac{x_j^* + x_{j-1}^*}{2} \right) - 1 \right|. \quad (20)$$

The estimates for the RND should not present spikes and non-differentiable points since it is difficult for meaningful economic interpretations. To avoid such scenarios, the function,

$$V(h) = \sum_{j=1}^m \left| f^{(2)}(x_j^*) - f^{(2)}(x_{j-1}^*) \right|, \quad (21)$$

is used to penalize the criterion function when extreme variations occur. Finally, entropy measures compensate for the tendency to define a flat density. Considering the same approach adopted for the area, a function,

$$\begin{aligned} H(h) \\ = - \sum_{j=1}^m (x_j^* - x_{j-1}^*) f^{(2)} \left(\frac{x_j^* + x_{j-1}^*}{2} \right) \log \left(f^{(2)} \left(\frac{x_j^* + x_{j-1}^*}{2} \right) \right). \end{aligned} \quad (22)$$

is defined as the approximation of an entropy measure.

The CV criterion function is tailored-made for this context. The bandwidth values h_c and h_p are defined by solving the problem,

$$\begin{aligned} \underset{h_c, h_p}{\text{minimize}} \quad & \frac{1}{n_c} \sum_{i=1}^{n_c} \frac{1}{k_{i,c}} \sum_{j=1}^{k_{i,c}} \left[\left(C_{i,j}^* - \hat{f}_{c(-k_{i,c})}(X_i) \right)^2 V_{c,p} + \frac{A_{c,p}}{H_{c,p}} \right] \\ & + \frac{1}{n_p} \sum_{i=1}^{n_p} \frac{1}{k_{i,p}} \sum_{j=1}^{k_{i,p}} \left[\left(P_{i,j}^* - \hat{f}_{p(-k_{i,p})}(X_i) \right)^2 V_{c,p} + \frac{A_{c,p}}{H_{c,p}} \right]. \end{aligned} \quad (23)$$

The standard optimization problem has been modified to apply the CV method to choose optimal bandwidths, including what can be interpreted as the addition of a regularization factor. Finding the optimal values of h_c and h_p that solve the

problem (23) represents a tremendous computational challenge, mainly when a grid-search approach is used. In a sequential framework, thinner grids are used for more accurate approximations; estimates are obtained by solving constrained quadratic programming problems.

3. Algorithm Design

We have utilized grid-based optimization methods to obtain the optimal values for equation (23), it is not possible to use gradient-based methods for the same since the assumptions of convexity and continuity cannot be postulated.

Algorithm 1 programs the minimization of an objective function through matrix-grid search. Given a matrix with the objective function values, lines are associated with the first parameter grid values, and columns are associated with the second. A correspondence is established between the function minimum value and grid indexes with each parameter.

Algorithm 1: Minimum grid-search.

```

Data:  $h_c = \{h_{c_1}, \dots, h_{c_n}\}; h_p = \{h_{p_1}, \dots, h_{p_m}\}; n; m; size = n \times m;$ 
 $cvvec = \{cv_1, \dots, cv_{size}\}$ 
Result:  $h_c^*; h_p^*$ 

 $min \leftarrow cvvec[1];$ 
 $row \leftarrow 1; col \leftarrow 1;$ 
for  $i \leftarrow 1$  to  $n$  do
  for  $j \leftarrow 1$  to  $m$  do
    if  $cvvec[(i-1) \times n + j] < min$  then
       $min \leftarrow cvvec[(i-1) \times n + j];$ 
       $row \leftarrow i; col \leftarrow j;$ 
    end
  end
end
 $h_c^* \leftarrow hc[row]; h_p^* \leftarrow hp[col];$ 

```

Algorithm 2 presents a sequential approach to populate the matrix associated with the grid search. The algorithm is based on a 3-layer loop, with an inner function *estimCVElements* which can take several seconds to compute.

Algorithm 2: CV matrix build sequential.

```

Data:  $h_c = \{h_{c_1}, \dots, h_{c_n}\}; h_p = \{h_{p_1}, \dots, h_{p_m}\}; strike = \{s_1, \dots, s_{n_u}\};$ 
 $size = n \times m$ 
Result:  $cvvec = \{cv_1, \dots, cv_{size}\}$ 

for  $k \leftarrow 1$  to  $n_u$  do
   $pos \leftarrow 0;$ 
  for  $i \leftarrow 1$  to  $n$  do
    for  $j \leftarrow 1$  to  $m$  do
       $pos \leftarrow pos + 1;$ 
       $[cv, var, area, entropy] \leftarrow$ 
         $estimCVElements(hc[i], hp[j], strike[k]);$ 
       $cvvec[pos] \leftarrow cvvec[pos] + cv \times var + \frac{1+|area-1|}{entropy};$ 
    end
  end
end

```

Algorithm 3 presents the code for the *estimCVElements* function. To define the output elements *cv*, *area*, *entropy*, and *var* may require several tens of thousands of constrained quadratic optimization problems to be solved.

Algorithm 3: estimCVElements.

Data: $x = \{x_1, \dots, x_m\}$; $\{callprice, putprice\}$; x_{ex}
Result: cv ; $area$; $entropy$; var

$ddf \leftarrow$ allocate vector density estimates dimension m ;
for $i \leftarrow 1$ **to** m **do**
 $ddf[i] \leftarrow estimxex(x[i], x_{ex})$;
end

$area \leftarrow areaEstim(x, ddf)$;
 $entropy \leftarrow entropyEstim(x, ddf)$;
 $var \leftarrow varEstim(ddf)$;

$[fcall, fput] \leftarrow estimxex(x_{ex}, x_{ex})$;

$cvcall \leftarrow 0.0$;
for $i \leftarrow 1$ **to** n_c **do**
 if $x[i] = x_{ex}$ **then**
 $cvcall \leftarrow cvcall + (callprice[i] - fcall)^2$;
 end
end

$cvput \leftarrow 0.0$;
for $i \leftarrow 1$ **to** n_p **do**
 if $x[i] = x_{ex}$ **then**
 $cvput \leftarrow cvput + (putprice[i] - fput)^2$;
 end
end

$cv \leftarrow cvcall + cvput$;

Algorithm 4 represents the CUDA/C++ version of the parallelization associated with the matrix-grid calculations. Comparing algorithm 4 with 2, we can see that the 3-layer loop is not present and the *estimCVElements* algorithm is evaluated in parallel. For the most complex problem in this experiment, $n \times m \times n_u = 5,898,240$ threads are run in parallel, each evaluating a unique instance of the function.

Algorithm 4: CV matrix build parallel.

Data: $h_c = \{h_{c_1}, \dots, h_{c_n}\}$; $h_p = \{h_{p_1}, \dots, h_{p_m}\}$; $strike = \{s_1, \dots, s_{n_u}\}$;
 $size = n \times m$;
Result: $cvvec = \{cv_1, \dots, cv_{size}\}$

/ allocate memory on the host */*
 $cvvecp \leftarrow$ allocate vector of dimension $n_u \times n_c \times n_p$;

/ running on the device */*
 $xid \leftarrow gridDim.x$;
 $yid \leftarrow gridDim.y$;
 $idx \leftarrow blockIdx.x$;
 $idy \leftarrow blockIdx.y$;
 $idz \leftarrow threadIdx.x$;
 $idt \leftarrow idz \times xid \times yid + xid \times idx + idy$;
 $[cv, var, area, entropy] \leftarrow estimCVElements(hc[idx], hp[idy], strike[idz])$;
 $cvvecp[idt] \leftarrow cv \times var + \frac{1+|area-1|}{entropy}$;

/ building cv matrix on the host */*
for $k \leftarrow 1$ **to** $size$ **do**
 $a \leftarrow 0.0$;
 for $i \leftarrow 1$ **to** n **do**
 $id \leftarrow (i - 1) \times size + k$;
 $a \leftarrow a + cvvecp[id]$;
 end
 $cvvec[k] \leftarrow a$;
end

4. Results

VIX Sequential

The following table tabulates the optimal kernel bandwidths h_c and h_p for Calls and Puts in VIX data for the sequential C++ algorithm, performed using Intel(R) Core(TM) i5-7300HQ CPU @ 2.50GHz. The grid [0.75 2.0] was assumed for h_c and h_p . The underlying [10.0 47.5] range within a grid of 128 values was considered. These values can be changed in the main.cpp file. Sample Size – Calls: 1465, Puts: 1196

Grid Dimensions	h_c	h_p	Elapsed Time (sec)
16x16	1.4167	1.0833	208.925
32x32	1.3952	1.0726	823.958
64x64	1.3849	1.0873	5178.250
128x128	1.3799	1.0846	13117.904

Table 1: Results for VIX index RND bandwidths (sequential)

VIX Parallel

The following table tabulates the optimal kernel bandwidths h_c and h_p for Calls and Puts in VIX data for the parallel CUDA C++ algorithm, performed using NVIDIA Tesla V100 – 16 GB GPU through the PACE-ICE Cluster at Georgia Tech. The grid [0.75 2.0] was assumed for h_c and h_p . The underlying [10.0 47.5] range within a grid of 128 values was considered. These values can be changed in the main.cu file. Sample Size – Calls: 1465, Puts: 1196

Grid Dimensions	h_c	h_p	Elapsed Time (sec)
16x16	1.4166	1.0833	1.78
32x32	1.3954	1.0727	6.06
64x64	1.3847	1.0872	22.12
128x128	1.3799	1.0847	87.48
256x256	1.3770	1.0788	353.57

Table 2: Results for VIX index RND bandwidths (parallel)

S&P 500 Sequential

The following table tabulates the optimal kernel bandwidths h_c and h_p for Calls and Puts in SP500 data for the sequential C++ algorithm, performed using Intel(R) Celeron(R) N4020 CPU @ 1.10GHz. The grid [0.25 1.25] was assumed for h_c and h_p . The underlying [24.0 28.0] range within a grid of 128 values was considered (though the algorithm terminated only for grid dimensions up to 32x32 within a day, only for which results are shown below). These values can similarly be changed in the main.cpp file. Sample Size – Calls: 9785, Puts: 13574

Grid Dimensions	h_c	h_p	Elapsed Time (sec)
16x16	0.5833	0.6500	12828.699
32x32	0.6048	0.6694	121803.956
64x64	0.6151	0.6786	488570.577

Table 3: Results for SP500 index RND bandwidths (sequential)

Running with grid dimensions of 64x64 on an Intel(R) Celeron(R) N4020 CPU @ 1.10GHz takes almost 6 days.

S&P 500 Parallel

The following table tabulates the optimal kernel bandwidths h_c and h_p for Calls and Puts in SP500 data for the parallel CUDA C++ algorithm. The grid [0.25 1.25] was assumed for h_c and h_p . The underlying [24.0 28.0] range within a grid of 256 values was considered. These values can similarly be changed in the main.cpp file. Sample Size – Calls: 9785, Puts: 13574

Grid Dimensions	h_c	h_p	Elapsed Time (sec)
16x16	0.5841	0.6512	16.42
32x32	0.6045	0.6713	45.60
64x64	0.6136	0.6799	218.52
128x128	0.6205	0.6820	758.27
256x256	0.6018	0.6710	3561.69

Table 4: Results for SP500 index RND bandwidths (parallel)

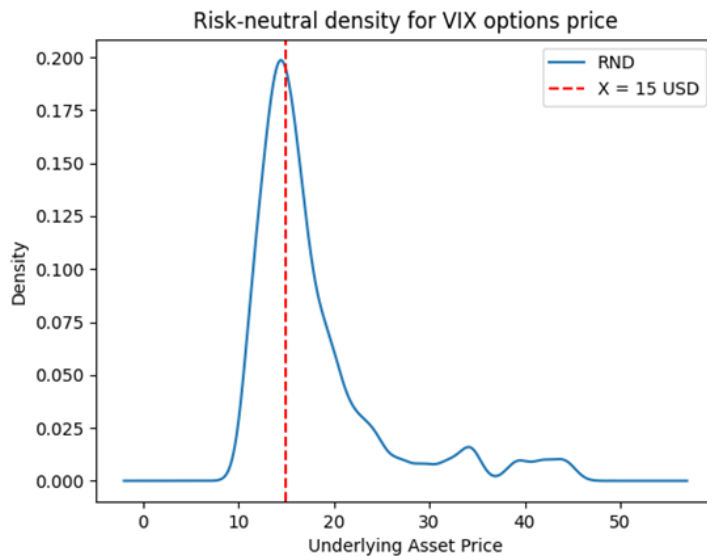


Figure 1: Risk-Neutral Density for VIX options price

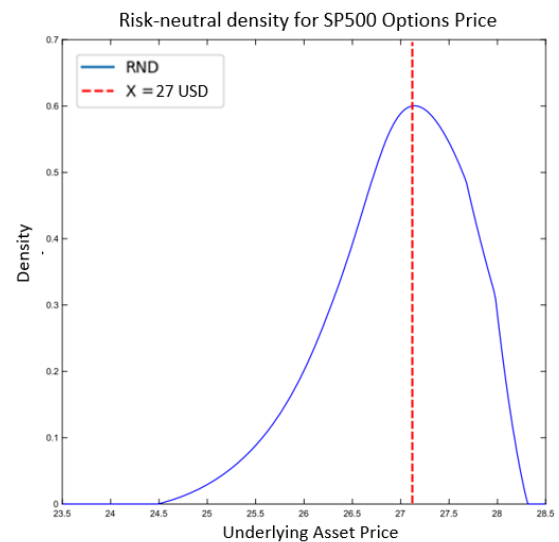


Figure 2: Risk-Neutral Density for SP500 index

Figure 1 represents the Risk-Neutral Density function for VIX options prices obtained from the h_c and h_p bandwidth values calculated using the CUDA/C++ code for 256 grid values. The RND function is generated from the geometric product of two Gaussian Kernels for Call and Put prices using the h_c and h_p bandwidth values respectively. We can see that the density peaks at around 15 USD, which indicates that the most probable underlying asset price for VIX index is 15 USD at the end of contract maturity. Based on this information, the portfolio manager can assess whether there is net positive risk in buying options contracts above this predicted market price.

Similarly, Figure 2 indicates the Risk-Neutral Density function for SP500 option prices that were similarly obtained from the h_c and h_p bandwidth values calculated using the CUDA/C++ code. We can see that the density peaks between 27 and 27.5 USD, corresponding to the most probable underlying asset price for SP500 index at maturity. Again, this information can inform decision making by portfolio managers by improving their understanding of risk at prices relative to this predicted market price.

5. Inferences and Conclusion

In this project, we have derived the optimal kernel bandwidths for RND options pricing through both sequential and parallel computing implementations. We have used a grid-based optimization method with a tailor-made cross-validation criterion function using sequential C++ codes run on local CPUs and CUDA/C++ parallel codes run on the NVIDIA Tesla V100 GPU through the PACE-ICE cluster with grid sizes ranging in 16x16, 32x32, 64x64, 128x128 and 256x256 dimensions.

For the VIX RND bandwidth values, we ran a sample size of 1465 Calls and 1196 Puts over 29 Strikes recorded between 16 to 20 April 2018 for contracts maturing on May 18, 2018. Since timing is imperative for making decisions regarding equity price expectations and assessing financial risks associated with options contracts, lower computational times are preferred, especially considering the Big Data required to generate the non-parametric RN densities for the options pricing function.

As we have seen from Table 1, the computational times associated with higher matrix grid dimensions are significant and exponentially increasing, and may even take a couple of days to run larger amounts of data. This opens up the possibility of the options pricing function losing relevance in real-time decision-making as the optimal values take longer and longer to estimate. Thus, it is imperative for a parallel computing alternative to be implemented so that computational times are shortened to within a couple of hours at higher grid dimensions.

We can see that for higher grid values, we have sequential processing times exceeding 120 hours which is mainly due to the large sample size and intensive computation associated with Equation 23. Comparing this with the parallel computing implementation shown in Table 2, we can see exponentially smaller (up to 150x) computational times which allows for real-time decision-making in portfolio optimization.

For the SP500 RND bandwidth values, we ran a sample size of 9785 Calls and 13574 Puts over 90 strikes, also recorded between 16 to 20 April 2018 for contracts maturing on May 18, 2018. As we can see in Table 2, computational times in sequential operations are significantly higher owing to the larger sample size, with higher grid dimensions (64x64 and above) taking several days to compute. In practical applications, these computational times may far exceed the affordable threshold for real-time decision-making in financial risk which, accounting for the time required for auxiliary decisions, may even exceed the time to maturity. This would render the computation irrelevant, which is why we require faster computational processes using parallel computing methods. However, with parallel computing implementation as seen in Table 4, we again see exponentially smaller (over 2,200x) computational times better allowing for real-time decision-making.

After calculating the optimal parameter values, we have visualized the risk-neutral density functions for both SP500 and VIX data and determined the expected asset value at the time of contract maturity. Depending on the Put price at the time of purchase, financial managers can assess the probability of net positive or negative returns derived from the Options contract at the time of maturity.

Future research in this area is in the domain of determining RND bandwidths through Multi-Layered Perceptron Network (MLP - Deep Learning) which involve training models with explicit options data through distributed GPU systems and does not depend on non-parametric kernel density estimation. The advantage of this approach is that it can learn complex patterns and dependencies in the RND without making any assumptions about the underlying distribution.

6. References

- [1] [D.T. Breeden, R.H. Litzenberger, Prices of state-contingent claims implicit in option prices, J. Bus. \(1978\) 621-651](#)
- [2] [R.W. Banz, M.H. Miller, Prices for state-contingent claims: some estimates and Applications. J. Bus \(1978\) 653-672](#)
- [3] [Ana M. Monteiro, Antonio A.F. Santos, Parallel computing in finance for estimating risk-neutral densities through options prices, Journal of Parallel and Distributed Computing, 2021](#)

Appendix

1. Main Function

```
#include <stdio.h>

#include <stdlib.h>

#include <math.h>

#include <time.h>

#include <string.h>

#include <cuda.h>

#include <cuda_runtime.h>

#include "rungpu.cuh"


int main(int argc, char *argv[])
{
    printf("Starting estimating ... \n");

    //

    double tau = 1.0/12.0;

    double r = 0.02;

    int ngrid = atoi(argv[1]);

    //

    run_in_gpu(0.75, 2.0, ngrid,

               0.75, 2.0, ngrid,

               10.0, 47.5, 128,

               tau, r);

    //

    //

    return 0;
}
```

2. GPU Run function

```
#ifndef rungpu_cuh
#define rungpu_cuh
//
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <time.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "io.cuh"
#include "kernelfunc.cuh"
//
//
//
__host__ __device__ void minMatrix(double *res, double *xxx, int nx, double *yyy, int
ny, double *mat)
{
    double min = mat[0];
    int row = 0;
    int col = 0;
    for(int i=0; i<nx; i++){
        for(int j=0; j<ny; j++){
            if(mat[i*nx + j] < min){
                min = mat[i*nx + j];
                row = i;
                col = j;
            }
        }
    }
}
```

```

    }

res[0] = xxx[row];
res[1] = yyy[col];
}

//

//

void run_in_gpu(double hmin,double hmax,int nhc,

                double hpmin,double hpmax,int nhp,

                double xmin,double xmax,int nx,

                double tau,double r)

{
int nc;

double *callprice = readArray2cuda("callprice.txt",&nc);
double *callstrike = readArray2cuda("callstrike.txt",&nc);
double *callopenint = readArray2cuda("callopenint.txt",&nc);
printf("Number observations (call) : %d\n",nc);

//

int np;

double *putprice = readArray2cuda("putprice.txt",&np);
double *putstrike = readArray2cuda("putstrike.txt",&np);
double *putopenint = readArray2cuda("putopenint.txt",&np);
printf("Number observations (put) : %d\n",np);

//

int nu;

double *strike = readArray2cuda("strike.txt",&nu);
printf("Number of strikes: %d\n",nu);

//

double time_spent = 0.0;

```

```

clock_t begin = clock();

//
double *hc = linspaceCuda(hcmin,hcmax,nhc);
double *hp = linspaceCuda(hpmin,hpmax,nhp);
double *xRange = linspaceCuda(xmin,xmax,nx);
unsigned int size = nhc*nhp*nu;
double *out = createCudaArray(size);
//
cudaDeviceSetLimit(cudaLimitMallocHeapSize,1024*1024*4000L);
kernel<<<dim3(nhc,nhp),nu>>>(out,
                                hc,nhc,hp,nhp,strike,nu,
                                callprice,callstrike,callopenint,nc,
                                putprice,putstrike,putopenint,np,
                                xRange,nx,
                                tau,r);

cudaDeviceSynchronize();
//
double *out_h = readArrayFromDevice(out,size);
double *cvvec = (double*)calloc(nhc*nhp,sizeof(double));
int ii = -1;
for(int k=0;k<nu;k++){
    int jj = -1;
    for(int i=0;i<nhc;i++){
        for(int j=0;j<nhp;j++){
            ii++;
            jj++;
            cvvec[jj] += out_h[ii];
        }
    }
}

```



```

}

double *hc_h = linspace(hcmin,hcmax,nhc);
double *hp_h = linspace(hpmin,hpmax,nhp);
//
double sol[2];
minMatrix(sol,hc_h,nhc,hp_h,nhp,cvvec);
printf("hcoptim: %.4f\n",sol[0]);
printf("hpoptim: %.4f\n",sol[1]);
//
writeMatrix(cvvec,nhc,nhp,"CVMatrix.txt");
//
clock_t end = clock();
time_spent = (double)(end - begin)/CLOCKS_PER_SEC;
printf("The elapsed time is %f seconds\n",time_spent);
//
cudaDeviceReset();
//
free(cvvec);
free(hc_h);
free(hp_h);
//
printf("Done ... \n");
// memory free zone
cudaFree(callprice);
cudaFree(callstrike);
cudaFree(callopenint);
cudaFree(putprice);
cudaFree(putstrike);
cudaFree(putopenint);

```

```

cudaFree (strike);

cudaFree (hc);

cudaFree (hp);

cudaFree (out);

cudaFree (xRange);

free (out_h);

}

////////////////////////////////////

#endif

```

3. Kernel function

```

#ifndef kernelfunc_cuh
#define kernelfunc_cuh

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "nprnd.cuh"

__global__ void kernel(double *out,
                        double *hc,int nhc,double *hp,int nhp,double *strike,int nu,
                        double *callprice,double *callstrike,double *callopenint,int
nc,
                        double *putprice,double *putstrike,double *putopenint,int np,
                        double *xRange,int mRange,
                        double tau,double r)
{

```

```

int ncc = gridDim.x;

int npp = gridDim.y;

int xi = blockIdx.x;

int yi = blockIdx.y;

int zi = threadIdx.x;

unsigned int size = nhp*nhc*nu;

int idx = zi*(ncc*npp - 1) + (ncc*xi + yi) + zi;

//

if(idx < size){

NPRND nprnd = NPRND(callprice,callstrike,callopenint,nc,

                    putprice,putstrike,putopenint,np,

                    r,tau,

                    xRange,mRange,

                    strike,nu);

out[idx] = nprnd.matCVElements(hc[xi],hp[yi],zi);

}

}

#endif

```

4. Graph Function

```

import numpy as np

import matplotlib.pyplot as plt

from scipy import stats

# Load call and put option prices from .txt files

option_prices_c = np.loadtxt('callprice.txt')

```

```

option_prices_p = np.loadtxt('putprice.txt')

# Define range of x values
x_range = np.linspace(-14, 45, len(option_prices_p))

# Define kernel function
def kernel_call(x, h):
    return np.exp(-0.5 * ((x - h) / h)**2) / np.sqrt(2 * np.pi * h**2)

def kernel_put(x, h):
    return np.exp(-0.5 * ((x + h) / h)**2) / np.sqrt(2 * np.pi * h**2)

# Define function to calculate risk-neutral density using Gaussian kernels
def risk_neutral_density(x, h_c, h_p):
    density_c = np.mean(kernel_call(x - option_prices_c, h_c)) / h_c
    density_p = np.mean(kernel_put(x - option_prices_p, h_p)) / h_p
    return density_c + density_p

# Set bandwidth parameters
h_c = 1.377
h_p = 1.078

# Calculate risk-neutral density
density = [risk_neutral_density(x, h_c, h_p) for x in x_range]

# peak_index=np.argmax(density)
# peak_x = x_range[peak_index]
# peak_y = density[peak_index]

# Plot density
plt.title('Risk-neutral density for VIX options price')

```

```
plt.plot(x_range+12, density, label='RND')  
plt.axvline(x=15, color='red', linestyle='--', label='X = 15 USD')  
plt.xlabel('Underlying Asset Price')  
plt.ylabel('Density')  
plt.legend()  
plt.show()
```